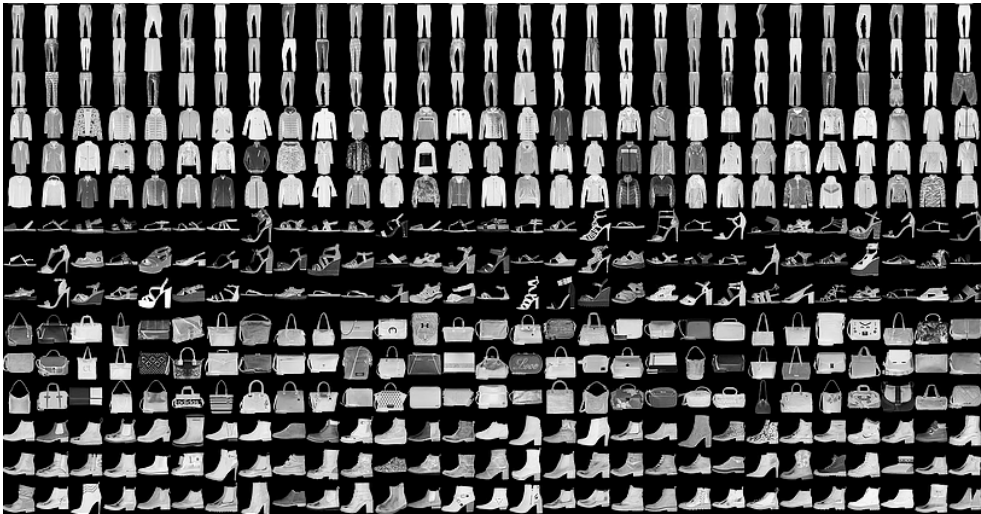


# Deep Learning (Homework 1)

Due date : 2022/4/1 23:55:00 (Hard Deadline)

## 1 Feedforward Neural Network (60%)

You are given a dataset sampled from the [Fashion-MNIST](#) dataset. This dataset contains 5 classes. In this exercise, you need to implement a **feedforward neural network (FNN)** model by yourself to recognize images, and use the backpropagation algorithm to update the parameters.



Label	Description
0	Trouser
1	Coat
2	Sandal
3	Bag
4	Ankle boot

Dataset description:

- [Training set](#) contained 30000 images with 6000 images collected for each individual class. [Test set](#) contained 5000 images with 1000 samples collected for each individual class.
- The images were 28 by 28 in size and were [flattened in row-major order into the shape of 784](#). The labels are integers that indicate the corresponding class. Details of these classes are provided in the above table.
- This dataset is given in the form of numpy array files (.npz) named “[train\\_x.npz](#)”, “[train\\_y.npz](#)”, “[test\\_x.npz](#)” and “[test\\_y.npz](#)”, where “[x](#)” indicates images and “[y](#)” indicates the corresponding labels.

Please follow the steps below to implement your program:

- Understand how the “forward pass” and “backward pass” in FNN work in accordance with the **backpropagation** algorithm.
- Both the training and test images need to be normalized (**divided by 255**).
- Use the **cross entropy** error function  $J(\mathbf{w}) = -\frac{1}{N} \sum_{n=1}^N \sum_{k=1}^K t_{nk} \log y_k(\mathbf{x}_n, \mathbf{w})$  as the objective function where  $t_{nk}$  is the target value,  $N$  is the number of samples in a batch and  $y_k(\mathbf{x}_n, \mathbf{w}_n)$  is the FNN output.

1. Design a FNN model architecture and use the file of the initial weights and biases “weights.npy”. Run the **backpropagation** algorithm and use the **mini-batch SGD** (stochastic gradient descent)  $\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla J(\mathbf{w}^{(\tau)})$  to optimize the parameters (the weights and biases), where  $\eta$  is the learning rate. **You should implement the FNN training under the following settings:**
  - number of layers: 3
  - number of neurons in each layer (in order): 2048, 512, 5
  - activation function for each layer (in order): relu, relu, softmax
  - number of training epochs: 20
  - learning rate: 0.001
  - batch size: 200
  - **important note:** For 1(a), **DO NOT RESHUFFLE THE DATA**. We had already shuffled the data for you. Reshuffling will make **your results differ from our ground-truth implementation**, and **any difference will result in reduction of your points**. On the same note, when splitting the samples into batches, split them in the given sample order.
  - (a) **Plot** the **learning curves** of  $J(\mathbf{w})$  and the **accuracy** of classification for every 50 iterations, with training data as well as test data, also, **show** the final loss and accuracy values. (20%)
  - (b) **Repeat 1(a)** by considering **zero initialization** for the model weights. And **do some discussion**. (8%)
2. Based on the model in 1, please **implement the dropout layers** and apply them **after the first two hidden layers**, i.e. the layers with 2048 and 512 neurons. The **dropout rate should be set as 0.2** for both layers. Note that the dropout operation **should only be applied in the training phase** and should be disabled in the test phase.
  - (a) **Train** the model using the same settings in 1 and **repeat 1(a)** (8%)
  - (b) Based on the experimental results, how the dropout layers affect the model performance and why? Please **do some discussion**. (8%)
3. In 1 and 2, the model is trained by **updating the weights  $\mathbf{w}$**  using their gradients, but in the process of backpropagation we can also obtain **the gradients of the input  $\mathbf{x}$** . What happens if we fix  $\mathbf{w}$  and update the input  $\mathbf{x}$  instead? **Using the trained model** in 1, please randomly generate a batch of (256) noise images as the initial input and perform backpropagation to **update these input  $\mathbf{x}$  while fixing  $\mathbf{w}$** . You should use the **learning rate 0.1 and update for 5 iterations**. For the labels, you can randomly assign integer values in the range  $[0, 4]$ .
  - (a) **Show** the **classification accuracy** of these noise images with respect to the random labels **before and after the updates**. (8%)
  - (b) Based on the result, **explain your findings**. What do these updated inputs tell us about the trained model? (8%)

**Note:**


- When coding in Python, be careful to assign the value to variable (**mutable vs immutable object**). Double check the dimensions of your matrices.
- Normally, when training a deep neural network, **you should shuffle the data for each epoch**, but **for convenience of grading we restrict the data order in 1(a)**.

## 2 Convolutional Neural Network (40%)

In this exercise, you will construct a convolutional neural network (CNN) for image recognition by using **Stanford Dogs Dataset**. This dataset consists of 1600 dog images from 8 categories.

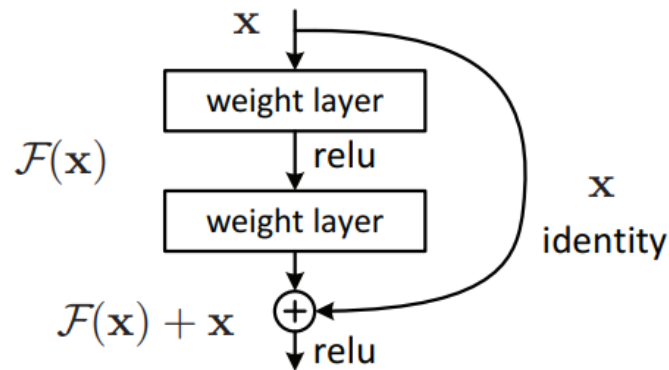


1. Please implement a CNN for image recognition by using **Stanford Dogs Dataset**, then **plot** the **learning curve** and the **accuracy rate** of training and test data. (20%)



**NOTE:** Figure above shows an example. The result might be different.

2. In order to deal with a real-world problem, we may stack some additional layers in the Deep Neural Network which results in the improved accuracy and performance. But it has been found that there is a maximum threshold for depth with the traditional convolutional neural network model. The problem of training a very deep network has been alleviated with the introduction of **ResNet** or residual network.



- (a) Construct a **ResNet** with residual blocks for image recognition and **plot** the [learning curve](#), [accuracy rate](#), try to stack more blocks as you can (ResNet-18 is recommended), you can refer to the [paper](#) for implementation. (15%)
- (b) Remove the identity mapping and repeat (a), then do some discussion on the results of (a) and (b). Please **describe** what you found. (5%)

**NOTE:** Please implement the model by yourself, directly load the pre-trained model from pytorch is not allowed.

### 3 Rule

- In your submission, you need to submit two files. And only the following file format is accepted:
  - **hw1\_<ProblemNumber>\_<StudentID>.ipynb** file which need to contain all the results, codes and reports for each exercise (e.g. **hw1\_2\_0123456.ipynb**).
- Implementation will be graded by
  - Completeness
  - Algorithm correctness
  - Description of model design
  - Discussion and analysis
- Only **Python** implementation is acceptable.
- For **problem 1**, any tools with **automatic differentiation** are **forbidden**, such as **Tensorflow**, **PyTorch**, **Keras**, etc. You should implement backpropagation algorithm **by yourself**.
- For **problem 2**, you should use **PyTorch** to implement the model, **other deep learning APIs are forbidden**.
- **DO NOT PLAGIARIZE**. (We will check program similarity score.)